

Yet Another CUDA FFT

Usage

Compiling the program

Type `make` to compile the program. Alternatively, type the following commands:

```
nvcc --compiler-options=-Wall -g -c argparse.c
nvcc --compiler-options=-Wall -g argparse.o HugoRiveraA3.cu -o fft -lm
```

The files `argparse.h` and `argparse.c` are used for command line argument parsing, thanks to the lightweight `argparse` library.

Usage

```
$ ./fft -h
Usage: fft [options]
```

Compute the FFT of a dataset with a given size, using a specified DFT algorithm.

```
-h, --help          show this help message and exit
```

Algorithm and data options

```
-a, --algorithm=<str>  algorithm for computing the DFT (dft|fft|gpu|fft_gpu|dft_gpu),
                        default is 'dft'
-f, --fill_with=<int>  fill data with this integer
-s, --no_samples       do not set first part of array to sample data
-N, --data_length=<int> data length
```

Benchmark options

```
-t, --measure_time=<int> measure runtime. runs algorithms <int> times. set to 0 if not needed.
-p, --no_print          do not print results
```

Measuring runtime

Runtime is easy to measure.

```
$ ./fft --measure_time=10 --no_print -N1024 -afft
```

Running Cooley-Tukey FFT with N=1024

```
0.00028737 (s)
0.00027086 (s)
0.00027070 (s)
0.00027063 (s)
0.00027062 (s)
0.00027062 (s)
0.00027062 (s)
0.00027062 (s)
0.00027062 (s)
0.00027062 (s)
0.00027062 (s)
```

```
$ ./fft --measure_time=10 --no_print -N1024 -afft_gpu
```

Running Cooley-Tukey FFT on GPU with N=1024

```
0.00054887 (s)
0.00044085 (s)
0.00044584 (s)
0.00042513 (s)
0.00042042 (s)
0.00041740 (s)
0.00041829 (s)
0.00041808 (s)
```

```
0.00041718 (s)
0.00041853 (s)
```

The FFT on the GPU only starts to outperform the FFT on the CPU on larger datasets.

```
$ ./fft --measure_time=10 --no_print -N65536 -afft
Running Cooley-Tukey FFT with N=65536
0.02675756 (s)
0.02649335 (s)
0.02648379 (s)
0.02648249 (s)
0.02648116 (s)
0.02648694 (s)
0.02650917 (s)
0.02648482 (s)
0.02648311 (s)
0.02648319 (s)
$ ./fft --measure_time=10 --no_print -N65536 -afft_gpu
Running Cooley-Tukey FFT on GPU with N=65536
0.00158091 (s)
0.00115752 (s)
0.00116558 (s)
0.00115046 (s)
0.00115190 (s)
0.00116676 (s)
0.00114784 (s)
0.00114956 (s)
0.00114897 (s)
0.00117116 (s)
```

Performance

In seconds

N	fft_gpu	fft	dft_gpu	dft
256	0.00041 ± 2.7e-05	6.99e-05 ± 9.4e-06	0.0004048 ± 1.7e-05	0.01285 ± 0.0014
512	0.00044 ± 4.2e-05	0.000137 ± 1.7e-07	0.0005946 ± 1.8e-05	0.04353 ± 0.0029
1024	0.00048 ± 3.1e-05	0.000277 ± 1.2e-05	0.00128 ± 2.3e-05	0.4002 ± 0.67
2048	0.00049 ± 2.7e-05	0.000468 ± 5.2e-06	0.004396 ± 0.00066	2.069 ± 1.0
4096	0.00047 ± 1.9e-05	0.00108 ± 2e-05	0.0155 ± 0.00091	
8192	0.00062 ± 3.7e-05	0.00211 ± 6.2e-05		
16384	0.00095 ± 2.7e-05	0.00454 ± 0.00017		
32768	0.00185 ± 0.00032	0.00924 ± 0.00066		
65536	0.00349 ± 0.00048	0.0187 ± 0.0033		
131072	0.00763 ± 0.0019	0.0308 ± 0.0026		
262144	0.0146 ± 0.0025	0.0621 ± 0.0026		
524288	0.0253 ± 0.0028	0.137 ± 0.002		

Speedup and Efficiency

The scripts `time.sh` and `plot.py` are used to gather and plot timing data from multiple runs.

Definition of the DFT

Let \mathbf{x} be an N dimensional complex vector. Then the DFT of \mathbf{x} is an N dimensional complex vector called \mathbf{Y} where each element of \mathbf{Y} is defined as follows:

$$Y[k] = \sum (x[n] * \exp(-2i * \pi * n * k / N)) \text{ where } n=0 \text{ to } N-1$$

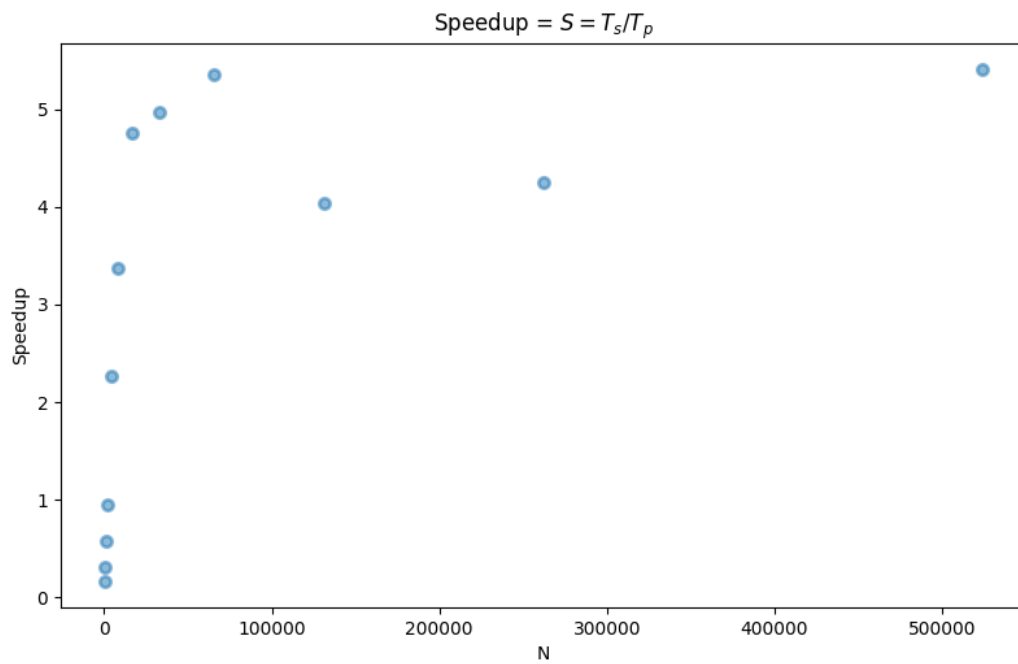


Figure 1: speedup plot

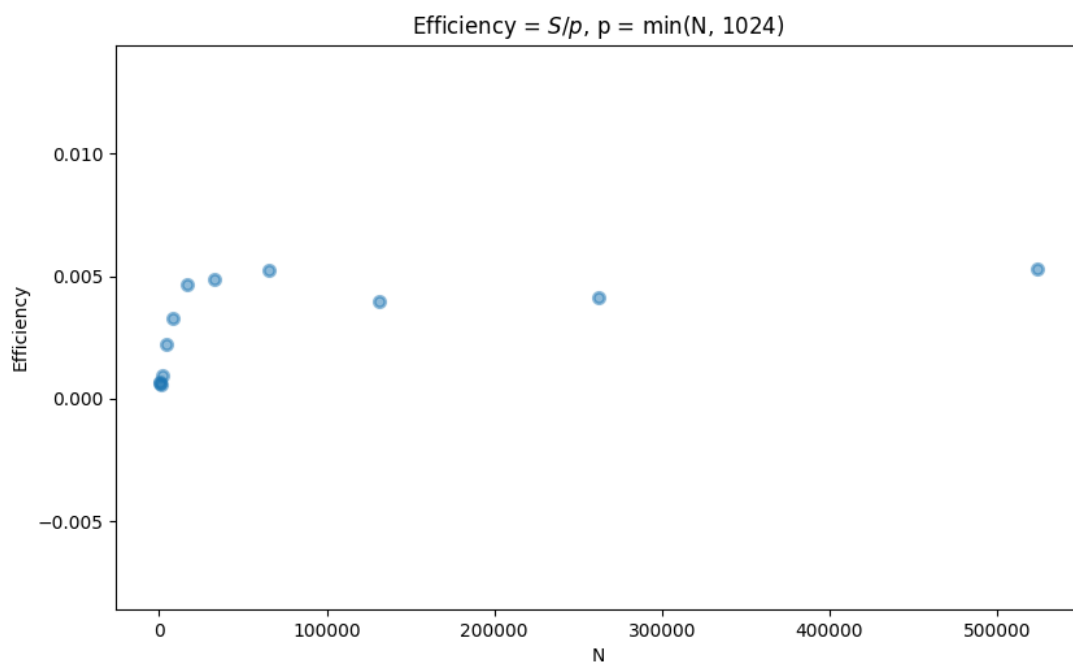


Figure 2: efficiency plot